

# Arrays and Strings

Prof. Darrell Long  
CSE 13S

# What is an array?

- It is a *homogeneous* collection of *elements*
  - All elements have the same type.
- Arrays can have one dimension
  - Often called a *vector*.
- Arrays can have two dimensions
  - Often called a *matrix*.
  - C treats a matrix as an array of vectors.
- Arrays can have even more dimensions
  - Often called a *tensor* by the machine learning community.
  - C treats these as arrays of arrays of ... some type.

# Arrays are ordered

- In mathematics we write a vector as  $\vec{a} = \langle a_1, \dots, a_n \rangle$ .
- In C, arrays start at zero:  $a[0], \dots, a[n-1]$ 
  - Why? It's easier, since  $a[0]$  is the base address of the array.
- $a[i]$  comes *before*  $a[i+1]$  in memory.
  - What about  $a[n]$ ?
    - It's past the end of the array, *don't do that!*
- $a[i]$  comes *after*  $a[i-1]$  in memory.
  - Does  $a[-2]$  make sense?
    - No, it's before the beginning of the array, *don't do that!*
    - Yes, but we need to talk about pointers, and some people use it, but do not be *that guy*.

# Declaring an array

---

*type name [ count ] = { initialization list }*

```
int a[] = {1, 2, 3, 4, 5, 6, 7, 8, 9, 10}; // The first 10 positive integers  
int b[10]; // Ten empty slots (usually zero, depending on where declared)  
float c[] = {3.1416, 2.7183, 0.57722, 1.6180}; // Four important constants
```

- If you have an initialization list, then you don't need a count.
- If you have a count, then you don't need an initialization list.
  - But you can have one, of course.
  - If the list is too short, then C will fill the rest with zero.

# In Memory

The assembly language shows us what is really happening.

```
.section      __DATA,__data
.globl _a          ## @a
.p2align    4
_a:
.long     1          ## 0x1
.long     2          ## 0x2
.long     3          ## 0x3
.long     4          ## 0x4
.long     5          ## 0x5
.long     6          ## 0x6
.long     7          ## 0x7
.long     8          ## 0x8
.long     9          ## 0x9
.long    10          ## 0xa

.globl _c          ## @c
.p2align    4
_c:
.long    1078530041    ## float 3.14159989
.long    1076754593    ## float 2.7183001
.long    1058260145    ## float 0.577220023
.long    1070537376    ## float 1.61800003

.comm    _b,40,4        ## @b
```

# Matrices

- In mathematics we write a matrix as  $m = \begin{bmatrix} m_{1,1} & m_{1,2} & m_{1,3} \\ m_{2,1} & m_{2,2} & m_{2,3} \\ m_{3,1} & m_{3,2} & m_{3,3} \end{bmatrix}$
- In **C**, we write: *type* `m[ 3 ][ 3 ]`;
  - The elements are `m[ 0 ][ 0 ]` to `m[ 2 ][ 2 ]`
  - **C** stores the array in row major order
    - That means one row in memory, then the next row, then the next row.
- FORTRAN is the only language that I know of that stores array in column major order.
  - That means one column, then the next, ...

## In Memory

- Memory is one dimensional.
- Rows are laid out sequentially.

```
int m[2][2] = {{1, 2}, {3, 4}};
```

```
.section      __DATA,__data
.globl _m          ## @m
.p2align    4
_m:
.long     1          ## 0x1
.long     2          ## 0x2
.long     3          ## 0x3
.long     4          ## 0x4
```

# Matrix Multiplication

---

```
void matMul(int n, int m, int p, float a[n][m], float b[m][p], float c[n][p]) {  
    for (int i = 0; i < n; i += 1) {  
        for (int j = 0; j < p; j += 1) {  
            c[i][j] = 0.0;  
            for (int k = 0; k < m; k += 1) {  
                c[i][j] += a[i][k] * b[k][j];  
            }  
        }  
    }  
    return;  
}
```

# How are we changing `C[ ][ ]?`

- Arrays are the exception to the rule that parameters in **C** are always passed by value.
- Why is this?
  - Arrays can be large, so you don't want to copy them onto the stack.
- Does this make sense?
  - Perhaps surprisingly, it makes perfect sense!
  - Arrays and pointers are intimately related.
- The name of the array is just a pointer to element *zero* of the array.

# & (address of) operator

---

& gives the memory location (address) of a variable.

```
#include <inttypes.h>
#include <stdio.h>

int main(void) {
    float x[2] = {0.1, 3.1415};

    uint64_t x0 = (uint64_t) &x[0];
    uint64_t x1 = (uint64_t) &x[1];

    printf("x[0] is at %" PRIu64 " and x[1] is at %" PRIu64 "\n", x0, x1);
    return 0;
}
```

x[0] is at 140732850817744 and x[1] is at 140732850817748

# sizeof operator

- The `sizeof` operator tells us the number of bytes used by a variable.
- Works on arrays, structures, unions, when the *compiler can* know how much memory is used.

```
#include <inttypes.h>
#include <stdio.h>

int main(void) {
    double x[] = {1.6180, 2.7183, 0.57722, 3.1416};

    uint64_t xs = sizeof(x);
    uint64_t x1 = sizeof(&x[1]);

    printf("&x takes %" PRIu64 " bytes and x takes %" PRIu64
           " bytes and x[1] takes %" PRIu64 " bytes\n",
           (uint64_t)sizeof(&x), xs, x1);
    return 0;
}
```

```
&x takes 8 bytes and x takes 32 bytes and x[1] takes 8 bytes
```



Is this what  
you  
expected?

```
#include <stdio.h>

int main(void) {
    char *s = "Hello world";
    char t[] = "Goodbye cruel world";

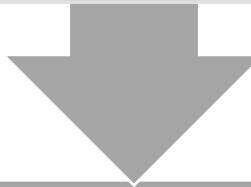
    printf("sizeof(s) = %lu, sizeof(*s) = %lu\n", sizeof(s), sizeof(*s));
    printf("sizeof(t) = %lu, sizeof(*t) = %lu\n", sizeof(t), sizeof(*t));
    return 0;
}
```

```
sizeof(s) = 8, sizeof(*s) = 1
sizeof(t) = 20, sizeof(*t) = 1
```

- `sizeof` when applied to an array gives the size of the array
- `sizeof` when applied to a pointer gives the size of the pointer

Arrays and pointers in C are related in fundamental and often confusing ways.

We will talk more about this when we deal with pointers.



For a single dimensional array,

$$a[i] == *(a + i)$$

`a` is the address of `a[0]`

`a[i]` is the array slot that is at  
`a + i * sizeof(a[0])`

Pointers automatically do the  
multiplication by `sizeof`

# Pointers and arrays

```
#include <stdio.h>
#include <stdlib.h>

int *newArray(int elements) { return (int *)calloc(elements, sizeof(int)); }

int main(void) {
    int *a;
    int n;
    do {
        printf("n > 1: ");
        scanf("%d", &n);
    } while (n < 2);
    a = newArray(n);
    for (int i = 0; i < n; i += 1) {
        a[i] = i * i;
    }
    for (int i = 0; i < n; i += 1) {
        printf("%d^2 = %d\n", i, a[i]);
    }
    free(a);
    return 0;
}
```

Arrays and points are equivalent in C

```
int **makeMatrix(int n, int m) {
    int **t = calloc(n, sizeof(int *)); // Allocate a column of row pointers
    for (int i = 0; i < n; i += 1) {
        t[i] = calloc(m, sizeof(int)); // Allocate each row
    }
    return t;
}
```

A matrix is an array of pointers

# What is a string?

- In **C**, a string is an array of characters that ends in '\0'
- It can be written as:
  - `char s[] = "Hello world!"`
  - `char *s = "Goodbye cruel world!"`
  - `char s[] = {'L', 'e', 'g', 'a', 'l', 0}`
- Most people use the `*s` version.
- A 100 character (99 printable) empty string is
  - `char s[100]`
- A string is just an array of characters that we have agreed is terminated by a null (zero) character.



## Assignment and Copying

```
char *s = "Wally Wonder Badger";
char *t;

t = s;
```

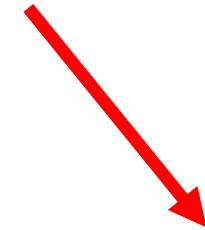
- What does that mean?
- It means that `s` and `t` point to (reference) the same memory location.
- It *does not* make a copy of `s`.

What is  
going on  
here?

```
#include <stdio.h>

int main(void) {
    int x[] = {1, 2, 3, 4};
    int *y;
    y = x;
    for (int i = 0; i < 4; i += 1) printf("%d ", x[i]); printf("\n");
    y[2] = 17;
    for (int i = 0; i < 4; i += 1) printf("%d ", x[i]); printf("\n");
    return 0;
}
```

- We'll use integers for simplicity.
- **y** is an *alias* for **x**
  - They reference the same memory.



```
1 2 3 4
1 2 17 4
```

# strcmp()

```
int strcmp(char s[], char t[]) {
    for (int i = 0; s[i] && t[i]; i += 1) {
        if (s[i] - t[i]) {
            return s[i] - t[i];
        }
    }
    return 0;
}
```

- As long as both characters are not *null*, subtract them.
  - < 0 means  $s < t$
  - > 0 means  $s > t$
  - = 0 means  $s == t$



strlen()

```
int strlen(char s[]) {
    int length = 0;
    while (s[length]) {
        length += 1;
    }
    return length;
}
```

- As long as `s[length]` is not *null*, move on to the next one.
- `length` is the number of non-null characters.

# strcpy()

```
void strcpy(char s[], char t[]) {
    for (int i = 0; t[i]; i += 1) {
        s[i] = t[i];
    }
    s[strlen(t)] = '\0'; // null terminate
    return;
}
```

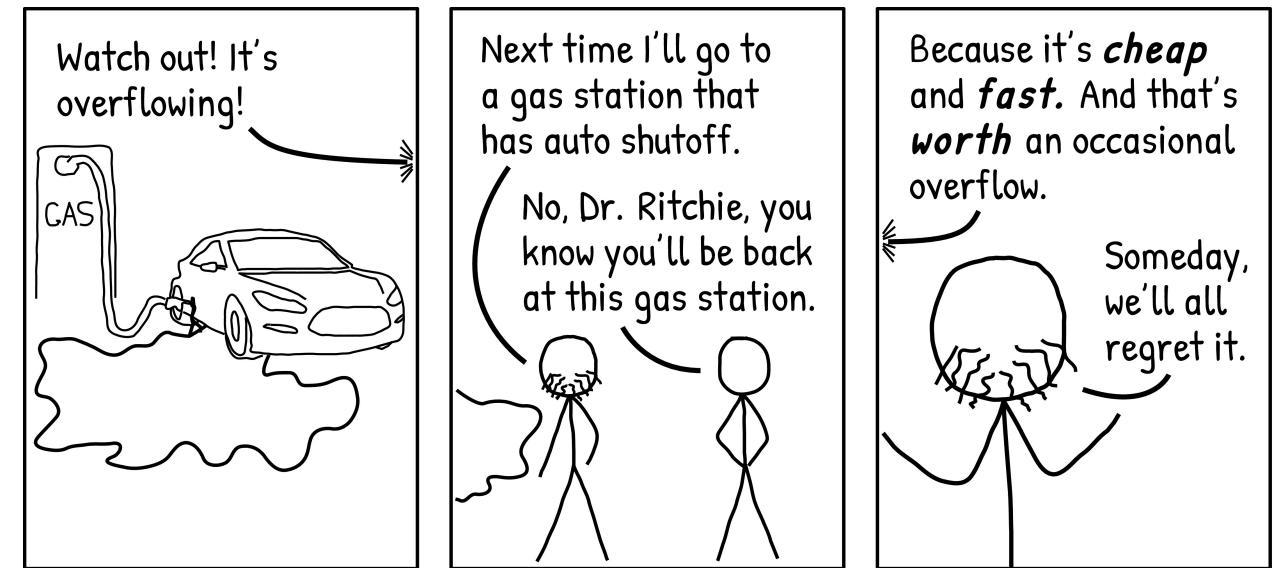
- As long as `t[i]` is not null, copy it to `s[i]`
- And be sure to tack a null on the end

You should never write code like that!

- Why not?
  - Did you check the array bounds?
- This is the source of most buffer-overflow attacks!



15 April 2021



Buffer Overflow.

# strcmp()

```
int strcmp(char s[], char t[], int n) {
    for (int i = 0; i < n && s[i] && t[i]; i += 1) {
        if (s[i] - t[i]) {
            return s[i] - t[i];
        }
    }
    return 0;
}
```

- We just have to add a quick check to make sure we haven't gone too far.
- It depends on you knowing how large the arrays are, of course.

# strncpy()

```
void strncpy(char s[], char t[], int n) {
    for (int i = 0; i < n && t[i]; i += 1) {
        s[i] = t[i];
    }
    s[strlen(t)] = '\0'; // null terminate
    return;
}
```

- Note that `n` must account for the null at the end.